

Homework 4

April 18, 2023

```
[1]: import statsmodels.api as sm
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
#1a
data_file = pd.read_csv("fifa22.csv")
print(data_file.head(5))
```

	name	rank	gender	wage_eur	log_wage	position	\
0	Lionel Andrés Messi Cuccittini	93	M	320000.0	12.676076	RW	
1	Lucia Roberta Tough Bronze	92	F	NaN	NaN	NaN	
2	Vivianne Miedema	92	F	NaN	NaN	NaN	
3	Wendéleine Thérèse Renard	92	F	NaN	NaN	NaN	
4	Robert Lewandowski	92	M	270000.0	12.506177	ST	

	nationality	club	league	preferred_foot	\
0	Argentina	Paris Saint-Germain	French Ligue 1	Left	
1	England	NaN	NaN	Right	
2	Netherlands	NaN	NaN	Right	
3	France	NaN	NaN	Right	
4	Poland	FC Bayern München	German 1. Bundesliga	Right	

	shooting	passing	dribbling	defending	attacking	skill	movement	power	\
0	92.0	91.0	95.0	26.333333	85.8	94.0	90.2	77.8	
1	61.0	70.0	81.0	89.000000	69.0	62.2	84.2	78.8	
2	93.0	75.0	88.0	25.000000	86.0	79.0	80.6	84.0	
3	70.0	62.0	73.0	91.333333	62.6	67.8	64.0	82.4	
4	92.0	79.0	86.0	32.000000	86.0	81.4	81.6	84.8	

	mentality	goalkeeping
0	73.833333	10.8
1	69.166667	12.6
2	70.833333	15.6
3	73.500000	12.8
4	80.666667	10.2

1c) The unit of analysis is a soccer player.

```
[2]: #1c
print("Number of observations:", data_file.shape[0])
print("Number of features:", data_file.shape[1])
```

```
Number of observations: 19630
Number of features: 20
```

```
[3]: #1d
print(data_file["gender"].value_counts())
```

```
M    19239
F      391
Name: gender, dtype: int64
```

1e) No, This dataset isn't representative about the real-world population of professional football players as this dataset only includes characters/players from FIFA 2022, and not all professional football/soccer players in the real world. Therefore, it probably won't be able to capture the representative of the real-world players.

```
[4]: #1f
data_file = data_file.dropna(subset=['passing'])
print(data_file.shape)
```

```
(17450, 20)
```

```
[5]: #2
x = data_file[['passing', 'attacking', 'defending', 'skill']]
y = data_file['rank']
x = sm.add_constant(x)
model = sm.OLS(y,x)
results = model.fit()
print(results.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          rank    R-squared:                0.705
Model:                  OLS    Adj. R-squared:            0.705
Method:                 Least Squares    F-statistic:        1.044e+04
Date:                   Tue, 18 Apr 2023    Prob (F-statistic):    0.00
Time:                   19:02:50    Log-Likelihood:        -47856.
No. Observations:       17450    AIC:                   9.572e+04
Df Residuals:           17445    BIC:                   9.576e+04
Df Model:                4
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]

const	25.3278	0.203	124.785	0.000	24.930	25.726
passing	-0.0247	0.010	-2.425	0.015	-0.045	-0.005
attacking	0.6109	0.006	94.005	0.000	0.598	0.624
defending	0.1719	0.002	84.413	0.000	0.168	0.176
skill	0.0066	0.009	0.730	0.465	-0.011	0.024

Omnibus:	171.799	Durbin-Watson:	1.342
Prob(Omnibus):	0.000	Jarque-Bera (JB):	178.339
Skew:	0.234	Prob(JB):	1.88e-39
Kurtosis:	3.163	Cond. No.	790.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

/opt/conda/envs/dsua-111/lib/python3.7/site-

packages/numpy/core/fromnumeric.py:2542: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.

return ptp(axis=axis, out=out, **kwargs)

2b) The R-squared value is a measure of the proportion of variance in rank by our features. It is 0.705
 2c) Only attacking and defending are significant at the 1% level as their p values are less than 0.01
 2d) A unit increase is associated with 0.0066 unit increase in ranking

3a)(a) Based on the statsmodels output from Q2, we can expect the features to do a relatively good job at predicting rank. This is because the multiple regression model has a reasonably high R-squared value, indicating that a significant proportion of the variation in the dependent variable (rank) is explained by the independent variables (passing, attacking, defending, and skill).

```
[6]: #3b
X = data_file[['passing', 'attacking', 'defending', 'skill']]
Y = data_file['rank']
print(X.head())
print(Y.head())
```

	passing	attacking	defending	skill
0	91.0	85.8	26.333333	94.0
1	70.0	69.0	89.000000	62.2
2	75.0	86.0	25.000000	79.0
3	62.0	62.6	91.333333	67.8
4	79.0	86.0	32.000000	81.4

0	93
1	92
2	92
3	92
4	92

Name: rank, dtype: int64

```
[7]: #3c
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳ random_state=123)
print(X_train.head())
```

	passing	attacking	defending	skill
17226	52.0	48.0	59.333333	53.2
13548	48.0	55.0	12.666667	54.0
17874	59.0	46.2	58.000000	57.8
19599	47.0	40.6	46.666667	40.0
15629	49.0	51.8	25.666667	49.6

```
[8]: #3d)
lr_model = LinearRegression().fit(X_train, y_train)
print('Intercept:', lr_model.intercept_)
print('Coefficients:', lr_model.coef_)
```


```
Intercept: 25.167733064621757
Coefficients: [-0.02444506  0.61230756  0.17314968  0.00612364]
```

3e) The coefficients change slightly from 0.6109 in question 2 to 0.612 in question 3

```
[9]: #3f
y_predicted = lr_model.predict(X_test)
print(y_predicted[:3])
```

```
[64.57617047 72.78035994 70.46341746]
```

```
[10]: #3g
plt.scatter(y_test, y_predicted)
plt.xlabel('Actual Rank')
plt.ylabel('Predicted Rank')
plt.show()
```



```
[11]: #3h
from sklearn import metrics
print('Root Mean Squared Error:', np.sqrt(metrics.
      ↪mean_squared_error(y_test, y_predicted)))
```

Root Mean Squared Error: 3.744562639987198

3h) The RMSE(roughly 3.75) is our measure of the average error and it indicates a deviation(error) average between the model and the data. 3i) Based on the analyses conducted above, it appears that the model performs reasonably well in predicting player rank. The multiple regression analysis in the previous question found that passing, attacking, defending, and skill were all significant predictors of rank, and the overall model had a relatively high R-squared value.

```
[12]: #4a
print(data_file['preferred_foot'].value_counts())
```

```
Right    13044
Left     4406
Name: preferred_foot, dtype: int64
```

```
[13]: #4b
right_foot_count = data_file['preferred_foot'].value_counts()['Right']
total_count = len(data_file)
percentage = right_foot_count / total_count * 100
print("Percentage of players who prefer their right foot:", percentage)
```

Percentage of players who prefer their right foot: 74.75071633237822

```
[14]: #4c
X = data_file[['shooting', 'passing', 'dribbling', 'defending', 'attacking',
               'skill', 'movement', 'power', 'mentality', 'goalkeeping']]
print(X.head())
```

	shooting	passing	dribbling	defending	attacking	skill	movement	power	\
0	92.0	91.0	95.0	26.333333	85.8	94.0	90.2	77.8	
1	61.0	70.0	81.0	89.000000	69.0	62.2	84.2	78.8	
2	93.0	75.0	88.0	25.000000	86.0	79.0	80.6	84.0	
3	70.0	62.0	73.0	91.333333	62.6	67.8	64.0	82.4	
4	92.0	79.0	86.0	32.000000	86.0	81.4	81.6	84.8	

	mentality	goalkeeping
0	73.833333	10.8
1	69.166667	12.6
2	70.833333	15.6
3	73.500000	12.8
4	80.666667	10.2

```
[15]: #4d
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
x_scaled = scaler.fit_transform(X)
print(x_scaled[:3])
```

```
[[ 2.7843116  3.29664165  3.31535783 -1.39304935  3.40016362  3.54858025
  2.77464012  1.94428215  2.18061369  0.2816757 ]
 [ 0.59771861  1.22971891  1.87671942  2.13166681  1.59341682  0.59820902
  2.07280881  2.06650141  1.62369703  1.48110007]
 [ 2.85484686  1.72184337  2.59603862 -1.46804331  3.42167251  2.15689571
  1.65171003  2.70204154  1.82259584  3.48014068]]
```

```
[16]: #4e
from sklearn.model_selection import train_test_split

y = data_file['preferred_foot']
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.3,
                                                    random_state=456)
print(x_train[:3])
```


```
[[ -2.01208594 -1.42775318 -1.51435684  0.44430269 -1.82649677 -1.57281886
  -0.96846017 -1.01342387 -1.55868391  0.54821445]
 [ -0.46031026  0.34389488  0.74636066  0.4255542  -0.27785665  0.61676481
   0.69254058 -0.6467661  -0.30562141 -0.25140179]
 [ 0.31557757  0.1470451  0.12980134 -0.88684012  0.08779448  0.00442362
```

```
-0.05607947 -0.40232758 -0.36529105 -1.85063428]]
```

```
[17]: #4f
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

errors = list()
accuracy = list()
for k in range(1, 31):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(x_train, y_train)
    pred_k = knn.predict(x_test)
    accuracy.append(metrics.accuracy_score(y_test, pred_k))

plt.plot(range(1, 31), accuracy, label="accuracy rate k value")
plt.xlabel("k")
plt.ylabel("Accuracy")
plt.show()
```



```
[ ]:
```

```
[18]: k = 5
knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(x_train, y_train)
test_preds = knn.predict(x_test)
```

```
print(test_preds[:3])
```

```
['Right' 'Right' 'Right']
```

```
[19]: from sklearn import metrics
confusion_matrix = metrics.confusion_matrix(y_test, test_preds)
classification_matrix = metrics.classification_report(y_test, test_preds)
print("Confusion Matrix:")
print(confusion_matrix)
print("Classification Matrix:")
print(classification_matrix)
```

Confusion Matrix:

```
[[ 206 1120]
 [ 406 3503]]
```

Classification Matrix:

	precision	recall	f1-score	support
Left	0.34	0.16	0.21	1326
Right	0.76	0.90	0.82	3909
accuracy			0.71	5235
macro avg	0.55	0.53	0.52	5235
weighted avg	0.65	0.71	0.67	5235

4h) Approximately, there were 1120 predicted left, however there were only 206 true lefts. So there were an additional 914 lefts predicted 4i) The recall for left is relatively low, suggesting a poor prediction model. 4j) The model did a poor job in predicting the left foot, however for the right foot the recall was high as 0.90. Overall, a poor job.

```
[20]: import pandas as pd
X_scaled = pd.DataFrame(x_scaled)
print(X_scaled.head())
```

	0	1	2	3	4	5	6 \
0	2.784312	3.296642	3.315358	-1.393049	3.400164	3.548580	2.774640
1	0.597719	1.229719	1.876719	2.131667	1.593417	0.598209	2.072809
2	2.854847	1.721843	2.596039	-1.468043	3.421673	2.156896	1.651710
3	1.232536	0.442320	1.054640	2.262906	0.905132	1.117771	-0.290023
4	2.784312	2.115543	2.390519	-1.074325	3.421673	2.379565	1.768682

	7	8	9
0	1.944282	2.180614	0.281676
1	2.066501	1.623697	1.481100
2	2.702042	1.822596	3.480141
3	2.506491	2.140834	1.614369
4	2.799817	2.996099	-0.118132


```
[21]: np.random.seed(2022)
sampled_data = X_scaled.sample(n=5000, random_state=2022)
print(sampled_data.head())
```

```

      0      1      2      3      4      5      6  \
291    1.373606  2.115543  1.465680  1.550464  1.830015  1.748668  0.037498
501    1.937889  0.934444  1.876719 -0.849343  1.959068  1.377552  2.049414
8871   1.020930 -0.246654 -0.795038 -0.736852  1.120221 -0.979033 -1.576714
12793  0.456648 -0.345079  0.129801 -1.580534  0.173830 -0.552250  1.090245
7256  -1.377269 -0.541929 -1.206077  0.763027 -0.600490 -1.591375 -0.430389

      7      8      9
291    1.944282  2.180614  0.948023
501    1.870951  1.404908  0.148406
8871   0.990972  0.310965  0.281676
12793  1.039860 -0.703419 -1.051018
7256   0.013218 -0.206172  0.814753
```

```
[22]: from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
inertias = []
silhouette_scores = []
k_values = range(2, 21)
for k in k_values:
    kmeans = KMeans(n_clusters=k, random_state=789)
    kmeans.fit(sampled_data)

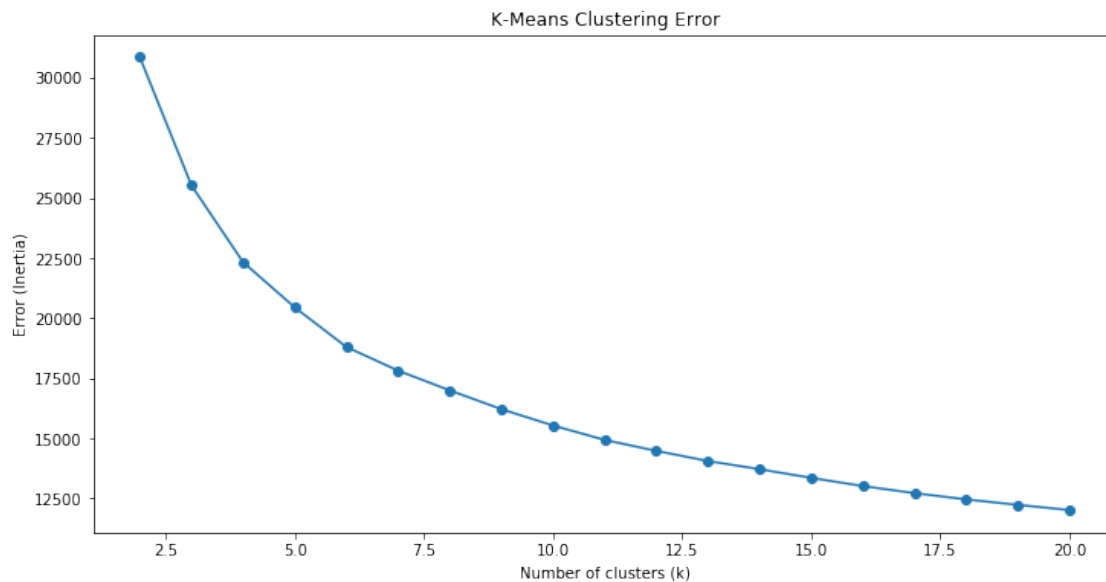
    inertia = kmeans.inertia_
    silhouette = silhouette_score(sampled_data, kmeans.labels_)

    inertias.append(inertia)
    silhouette_scores.append(silhouette)
print("error:", inertias)
```

```
error: [30847.126857997708, 25514.51134236195, 22325.868999171504,
20446.193863352895, 18799.73556250355, 17818.103488273115, 16997.425516956362,
16215.351018257887, 15536.930308302319, 14936.765072641238, 14481.726819497208,
14059.160188748174, 13719.567782417698, 13359.277645318618, 13018.39026372852,
12724.739592119795, 12464.422720758232, 12236.010193739769, 12022.748945684381]
```

```
[23]: plt.figure(figsize=(12,6))
plt.plot(range(2, len(inertias)+2), inertias, marker='o')
plt.title('K-Means Clustering Error')
plt.xlabel('Number of clusters (k)')
plt.ylabel('Error (Inertia)')
```

```
plt.show()
```



```
[27]: import sys
      !{sys.executable} -m pip install kneed
      import kneed
      from kneed import KneeLocator


      kl = KneeLocator(range(2, 21), inertias, curve='convex', direction='decreasing')
      elbow = kl.elbow

      print("The suggested elbow value is:",elbow)
```

```
Requirement already satisfied: kneed in
/opt/conda/envs/dsua-111/lib/python3.7/site-packages (0.7.0)
Requirement already satisfied: matplotlib in
/opt/conda/envs/dsua-111/lib/python3.7/site-packages (from kneed) (3.1.2)
Requirement already satisfied: numpy>=1.14.2 in
/opt/conda/envs/dsua-111/lib/python3.7/site-packages (from kneed) (1.18.1)
Requirement already satisfied: scipy in
/opt/conda/envs/dsua-111/lib/python3.7/site-packages (from kneed) (1.2.1)
Requirement already satisfied: python-dateutil>=2.1 in
/opt/conda/envs/dsua-111/lib/python3.7/site-packages (from matplotlib->kneed)
(2.8.1)
Requirement already satisfied: cycler>=0.10 in
/opt/conda/envs/dsua-111/lib/python3.7/site-packages (from matplotlib->kneed)
(0.10.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in
```

```
/opt/conda/envs/dsua-111/lib/python3.7/site-packages (from matplotlib->kneed)
(2.4.6)
Requirement already satisfied: kiwisolver>=1.0.1 in
/opt/conda/envs/dsua-111/lib/python3.7/site-packages (from matplotlib->kneed)
(1.1.0)
Requirement already satisfied: six>=1.5 in
/opt/conda/envs/dsua-111/lib/python3.7/site-packages (from python-
dateutil>=2.1->matplotlib->kneed) (1.13.0)
Requirement already satisfied: setuptools in
/opt/conda/envs/dsua-111/lib/python3.7/site-packages (from
kiwisolver>=1.0.1->matplotlib->kneed) (44.0.0.post20200106)
The suggested elbow value is: 6
```

```
[30]: plt.plot(range(2, len(silhouette_scores)+2), silhouette_scores)
plt.xlabel('Number of clusters')
plt.ylabel('Silhouette Score')
plt.show()
```



```
[41]: from sklearn.cluster import KMeans
k = 4

kmeans = KMeans(n_clusters=k)

kmeans.fit(x_scaled)
labels = kmeans.labels_
```

```
X['cluster_label'] = labels
```

```
print(X.head())
```

	shooting	passing	dribbling	defending	attacking	skill	movement	power	\
0	92.0	91.0	95.0	26.333333	85.8	94.0	90.2	77.8	
1	61.0	70.0	81.0	89.000000	69.0	62.2	84.2	78.8	
2	93.0	75.0	88.0	25.000000	86.0	79.0	80.6	84.0	
3	70.0	62.0	73.0	91.333333	62.6	67.8	64.0	82.4	
4	92.0	79.0	86.0	32.000000	86.0	81.4	81.6	84.8	

	mentality	goalkeeping	cluster_label
0	73.833333	10.8	2
1	69.166667	12.6	2
2	70.833333	15.6	2
3	73.500000	12.8	2
4	80.666667	10.2	2

/opt/conda/envs/dsua-111/lib/python3.7/site-packages/ipykernel_launcher.py:9:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.


Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
if __name__ == '__main__':
```

```
[46]: colors={
    0: 'C1',
    1: 'C0',
    2: 'C4',
    3: 'C3'
}
plt.rc('figure',figsize=(10,6))
for i in range(4):
    subset = X[X['cluster_label']==i]
    plt.scatter(subset['attacking'],subset['defending'],label='Cluster{}'.format(i),color=colors[i])

plt.xlabel('Attacking Score')
plt.ylabel('Defending Score')
plt.title('Player Clusters Based on Attacking and Defending Scores')
plt.show()
```



5i)Based on the analysis, clustering seems to be a meaningful technique for this data. The elbow plot suggested that a reasonable value for k would be 4 and the Silhouette Score plot shows that a k value of 4 has a relatively high Silhouette Score. 5j)I would have liked to run more clustering algorithms to see if they provided better results. This would helped in visualizing data in a better way.